

VetoPoint_{AI}

Final Business Plan

Competition-ready revised edition

Company slogan

Others secure the prompt. VetoPoint AI secures the action.

Core message

Existing security products protect what employees tell AI. VetoPoint protects organisations from what AI decides to do.

Team Trinorte

Company: VetoPoint AI | Product: VetoPoint Decision Gate | A Coruña, Galicia

Prepared for the AI Company Formation Competition

Contents

- 1 Competition Identity and Requirements Snapshot
- 2 Executive Summary
- 3 Mission, Vision, Values and Main Purpose
- 4 Team Structure and Roles
- 5 The Problem
- 6 The AI Solution
- 7 Role of AI and Measurable Improvement
- 8 Product Components
- 9 Target Customers and Buyers
- 10 Value Proposition
- 11 Company and Brand Design
- 12 AI Website and Landing Page Plan
- 13 One-Minute Promotional Video
- 14 Business and Revenue Model
- 15 Pricing Model
- 16 Market Opportunity
- 17 Go-to-Market, Marketing and Main Partners
- 18 Technology Model
- 19 Responsible AI Framework
- 20 Social, Ethical and Environmental Value
- 21 Development Roadmap
- 22 Illustrative Financial Projection
- 23 Main Costs
- 24 Investment and Support Required
- 25 Key Risks and Mitigation
- 26 Success Indicators
- 27 Five-Minute Investor Pitch Structure
- 28 AI Knowledge Quiz Readiness
- 29 Competition Requirements Compliance Matrix

1. Competition Identity and Requirements Snapshot

A concise identity block for the competition submission.

Identity element	Final wording
Company name	VetoPoint AI
Product name	VetoPoint Decision Gate
Proposed team name	Trinorte
Slogan	Others secure the prompt. VetoPoint AI secures the action.
One-sentence company description	VetoPoint AI is the enterprise control layer that evaluates, approves, escalates or vetoes consequential actions proposed by AI agents before execution.
Main purpose	Prevent harmful, unauthorised or non-compliant AI actions while enabling organisations to adopt autonomous AI safely and efficiently.
Primary customers	Banks, insurers, healthcare organisations, governments, technology firms and other regulated enterprises.
Primary call to action	Request a paid enterprise pilot.

Competition theme fit

VetoPoint is an AI-enabled company designed to solve a real and rapidly emerging business and societal problem: organisations need a safe, accountable way to govern actions taken by autonomous AI agents.

2. Executive Summary

VetoPoint AI is an enterprise AI-governance and action-control company that protects organisations from dangerous, unauthorised or non-compliant actions proposed by autonomous AI agents.

As AI evolves from answering questions to performing tasks, organisations face a new category of risk. AI agents may transfer money, send confidential information, modify records, approve transactions, change prices, sign contracts or delete data. Existing cybersecurity products largely focus on information entering AI systems. VetoPoint focuses on the actions coming out.

VetoPoint creates a control layer between an AI agent and the systems it can affect. Before a consequential action is executed, the platform:

- Identifies the AI agent and requested action
- Checks organisational policies and the agent's authority
- Calculates financial, legal, privacy, security and reputational risk
- Approves safe actions, escalates uncertain actions or vetoes prohibited actions
- Requests human approval when required
- Records the decision history for audit and compliance

VetoPoint will generate revenue through annual enterprise subscriptions, paid pilots, implementation services, premium compliance modules, private deployments and usage-based fees.

Investment proposition

VetoPoint seeks €750,000 to build a production-ready product, complete cybersecurity testing and launch three paid enterprise pilots.

3. Mission, Vision, Values and Main Purpose

Mission

Enable organisations to deploy autonomous AI safely by ensuring every consequential action is governed by policy, assessed for risk and subject to the appropriate level of human authority.

Vision

A future in which AI agents can act at enterprise speed without weakening human accountability, organisational trust or control.

Main purpose

Create the independent control point between AI intention and real-world execution, preventing harmful actions while allowing safe automation to proceed.

Core values

Value	Meaning
Human authority	People retain final responsibility for consequential decisions.
Safety by design	Controls are built into the action workflow, not added after an incident.
Transparency	Every recommendation, rule and final decision can be explained.
Accountability	Every action has an owner, a decision history and an auditable outcome.
Privacy and data minimisation	Only the information needed to assess an action is processed and retained.
Fairness	Policies and risk models are reviewed to avoid unjustified or discriminatory outcomes.
Controlled innovation	VetoPoint helps organisations adopt AI confidently rather than stopping useful innovation.

4. Team Structure and Roles

Proposed team name

Trinorte

Trinorte is a three-member founding team based in A Coruña, Galicia. For the competition, the members take complementary roles covering strategy, AI product development, and market communication.

Team member	Competition role	Primary responsibilities	Distinct contribution
Rubina Danilova	Marketing and Presentation Lead	Brand; website; promotional video; customer messaging; investor-style pitch delivery	Communication, visual identity and market engagement
Brais Muiño Otero	AI and Product Lead	Prototype; AI workflow; rules engine; technical architecture; product demonstration	AI-product design and technical implementation
Adrián Garrote Camoira	Team Leader and Business Strategist	Problem definition; business model; governance; customer discovery; partnerships; final pitch narrative	Business strategy, responsible AI and enterprise value

Team diversity: Trinorte combines three complementary perspectives - business strategy and governance, AI product development, and marketing and communication - together with local and international experience.

5. The Problem

AI agents are increasingly capable of taking real actions, not merely producing text. They can transfer funds, communicate with customers, approve invoices, modify records, change prices, access databases, create contracts, delete information and initiate purchases.

Central problem

Organisations are deploying increasingly autonomous AI systems without an independent control point that verifies whether their proposed actions are safe, authorised and compliant.

Who faces the problem?

- Regulated organisations using or preparing to use AI agents

- Executives responsible for risk and financial control
- Compliance, privacy and cybersecurity teams
- Employees asked to supervise automated workflows
- Customers and citizens affected by AI-generated actions

Why the problem matters

Financial risk: An AI agent may approve a payment to an unverified supplier or process an unusually large transaction.

Privacy risk: An agent may disclose personal, medical or financial information to an unauthorised recipient.

Legal and compliance risk: An action may violate organisational policy, contractual obligations or regulation.

Cybersecurity risk: A manipulated or compromised agent may access, alter or delete critical information.

Reputational risk: Incorrect automated decisions or communications may damage customer trust.

Accountability gap: The organisation may be unable to explain who authorised an action, which rule applied or why it was permitted.

6. The AI Solution

VetoPoint is an enterprise control platform positioned between AI agents and organisational systems. Every consequential AI action passes through VetoPoint before execution.

APPROVE	ESCALATE	VETO
The action complies with policy and falls within the agent's authorised limits.	The action requires human review because risk or uncertainty exceeds a defined threshold.	The action violates a non-negotiable rule or presents unacceptable risk.

Demonstration scenario

FinanceBot proposes: Transfer €48,000 to a newly registered supplier. No invoice is attached.

- Payment exceeds the automatic approval limit
- Supplier is not verified
- Required invoice is missing
- Transaction requires senior approval

VetoPoint result

Risk score: 92/100 | Risk level: Critical | Decision: Block pending review | Required approvers: CFO and Compliance Officer | Audit status: Recorded

7. Role of AI and Measurable Improvement

The role of AI in VetoPoint

- Interprets the context and intent of the proposed action
- Detects combinations of risk that fixed rules may miss
- Identifies unusual patterns and relevant risk factors
- Produces a structured risk score and explanation
- Recommends a safer alternative or appropriate approver
- Supports continuous improvement by analysing anonymised decision outcomes

Important control principle

AI provides contextual analysis and recommendations. It cannot override hard organisational rules or autonomously approve critical actions.

How the solution improves business performance

Dimension	Current challenge	VetoPoint improvement
Speed	Manual review of every automated action is slow.	Automatically approves low-risk actions and routes only exceptions to people.
Quality	Different teams may apply policies inconsistently.	Applies the same rules and structured risk criteria to every action.
Cost	Compliance teams spend time on repetitive checks and incident investigation.	Reduces routine review and creates audit evidence automatically.
Access	Safe autonomous AI may be limited to organisations with large governance teams.	Provides configurable policy templates and approval workflows.
User experience	Approvers may receive incomplete or unclear requests.	Presents a concise explanation, relevant policy and recommended next step.
Risk reduction	Problems are often discovered after execution.	Intervenes before money, data or systems are affected.

8. Product Components

8.1 VetoPoint Decision Gate

Evaluates proposed AI actions in real time using agent identity, action type, financial amount, target system, recipient, data sensitivity, records affected, documents, justification and applicable policies.

8.2 Policy Engine

Allows organisations to define deterministic rules. Hard rules always override AI recommendations.

8.3 Risk Analysis Engine

Assesses contextual financial, privacy, legal, fraud, cybersecurity, operational and reputational risks.

8.4 Human Approval Centre

Allows authorised staff to review, approve, reject, request information, add justification or reassign an approver.

8.5 Audit Trail

Records the proposed action, identified risks, applied policies, human decisions, reasons and final outcome.

8.6 Executive Dashboard

Shows actions analysed, high-risk actions prevented, approvals pending, financial value protected, common violations and risk trends.

8.7 Policy Builder

Provides a visual tool for configuring conditions, thresholds, required actions and approvers.

8.8 Agent Identity and Authority

Confirms which AI agent is acting, what systems it may access and what limits apply to its role.

Example policy

When a payment exceeds €25,000 and the recipient is a new supplier, require CFO and Compliance approval.

9. Target Customers and Buyers

Primary customer organisations

- Banks and payment companies
- Insurance companies
- Healthcare and pharmaceutical organisations
- Government agencies
- Technology companies
- Consulting and audit firms
- Large multinational enterprises

Initial market-entry profile

- Already uses AI agents or automated AI workflows
- Operates in a regulated or high-risk sector
- Processes high-value transactions or sensitive information
- Has dedicated compliance, risk or cybersecurity functions
- Requires documented human oversight and audit evidence

Primary buyers and decision-makers

- Chief Information Security Officer
- Chief Risk Officer
- Chief Compliance Officer
- Chief Information Officer
- Head of AI Governance
- Data Protection Officer
- Chief Financial Officer
- Digital Transformation Director

10. Value Proposition

For Executive leaders

- Safer adoption of autonomous AI
- Reduced financial and reputational exposure
- Clear accountability and organisation-wide visibility

For Compliance teams

- Enforceable policies
- Searchable audit evidence
- Consistent approval procedures

For Cybersecurity teams

- Control over agent access and execution
- Detection of suspicious actions
- Centralised governance across AI tools

For Employees and approvers

- Clear guidance
- Protection from accidental violations

- Concise explanations and safer alternatives

For Customers and citizens

- Greater protection from automated mistakes
- More transparent and accountable decisions

For Investors

- High-value enterprise market
- Recurring subscription revenue
- Potential to become control infrastructure for autonomous AI

11. Company and Brand Design

Element	Final wording
Company name	VetoPoint AI
Team name	Trinorte
Slogan	Others secure the prompt. VetoPoint AI secures the action.
Product line	Before AI acts, it passes VetoPoint.
One-sentence description	VetoPoint AI is the enterprise control layer that evaluates, approves, escalates or vetoes consequential actions proposed by AI agents before execution.

Brand colours

Governance Navy #14264C	Action Blue #2E9BEA	Signal Coral #F2503F	Cloud Gray #E9EAED	Trust White #FFFFFF
----------------------------	------------------------	-------------------------	-----------------------	------------------------

Logo meaning: the deep-navy V represents VetoPoint and organisational authority. The coral circle is the decisive control point, surrounded by a white boundary that signals governed review. The blue path rises through the checkpoint and becomes an arrow, representing an approved AI action proceeding safely. The visual language avoids generic robots, locks and shields.

12. AI Website and Landing Page Plan

Hero headline Before AI acts, it passes VetoPoint.

Hero subheading: Evaluate risk. Enforce policy. Approve, escalate or stop every consequential AI action before execution.

Primary calls to action: Test the Decision Gate | Request an Enterprise Pilot

Website section	Required content
Company introduction	Briefly explain VetoPoint as the control layer for autonomous AI actions.
Problem statement	Show how AI agents can transfer money, send data, change records and execute decisions before organisations can intervene.
AI solution	Explain the approve-escalate-veto model and the combination of policy rules, AI risk analysis and human authority.
How it works	Visual workflow from AI proposal to policy check, risk analysis, approval and audit record.
Target customers	Banks, insurers, healthcare, government and large enterprises.
Key benefits	Risk prevention, faster approvals, consistent policy enforcement and complete auditability.

Interactive demonstration	Three scenarios: risky supplier payment, customer-data transfer and low-risk internal analysis.
Responsible AI	Privacy, fairness, transparency, human oversight, contestability and responsible monitoring.
Team members	Three named members with roles and short biographies.
Contact / CTA	Pilot request form, business email and optional meeting-booking link.
Promotional video	Embed the one-minute video near the solution or CTA section.

Website and contact: Landing-page prototype prepared by Trinorte | hello@vetopoint.ai | Request-demo form included; add the final live URL before submission.

13. One-Minute Promotional Video

Timing	Scene	Visual direction	Voice-over / text
0-8 sec	The risk	Finance AI prepares a €48,000 payment to a new supplier without an invoice.	AI agents are beginning to act. Who checks their decisions?
8-16 sec	The problem	Fast examples: money transfer, data sharing, contract approval, record change and data deletion.	AI can act in seconds - before people see the risk.
16-29 sec	VetoPoint intervenes	The payment pauses. Risk score 92/100; new supplier; missing invoice; CFO approval required.	Action paused before execution.
29-41 sec	How it works	AI proposes -> policies checked -> risk analysed -> human approval. Show Approve, Escalate, Veto.	Every consequential action reaches a controlled decision point.
41-51 sec	Business value	Executive dashboard with clearly labelled demo data and regulated-sector examples.	Faster safe automation. Lower risk. Complete accountability.
51-60 sec	Closing	Logo and final call to action.	Others secure the prompt. VetoPoint AI secures the action. Request an enterprise pilot.

Why users or investors should support the idea

- AI-agent adoption creates a new action-governance problem that existing prompt-security tools do not fully solve.
- The product addresses direct financial, privacy, legal and operational risks.
- The business model combines high-value recurring subscriptions with implementation and premium modules.
- The platform enables useful AI automation while preserving human authority and accountability.

14. Business and Revenue Model

Revenue formula

Annual revenue = platform subscription + implementation + protected agents + action volume + premium modules + professional services

Who pays?

VetoPoint sells to organisations rather than individual employees. Primary paying customers are regulated enterprises, public-sector organisations and large companies deploying AI agents.

How the company earns income

- Annual enterprise software subscriptions

- Three-month paid pilots
- Initial configuration and integration fees
- Usage charges for additional protected actions
- Premium compliance, fraud, analytics and identity modules
- Private-cloud or on-premises deployment
- Training, security assessment and custom policy services

Revenue quality objective

At least 75% of mature revenue should come from recurring subscriptions and modules, reducing dependence on one-time consulting work.

15. Pricing Model

Paid pilot: VetoPoint Pilot - €25,000

Three-month controlled deployment for one department, including up to three AI agents, three approval workflows, risk analysis, policy enforcement, audit trail, dashboard and an evaluation report.

Annual subscription plans

Plan	Annual price	Typical customer
VetoPoint Core	€40,000	Small regulated organisation or one department
VetoPoint Business	€95,000	Medium organisation with several AI agents
VetoPoint Enterprise	€250,000+	Banks, governments and large enterprises

Implementation and premium revenue

- Implementation: €15,000 for small deployments; €30,000-€50,000 for medium deployments; €75,000-€150,000 for enterprise deployments.
- Premium modules: financial transaction control, privacy governance, compliance reporting, advanced fraud detection, agent identity verification, executive analytics and private deployment.
- Additional services: custom policy design, responsible-AI audits, security assessments, enterprise training and premium support.

16. Market Opportunity

VetoPoint targets the emerging enterprise need to govern AI systems that do more than generate content. The market opportunity grows as organisations connect AI agents to payments, databases, customer communication, workflow tools and operational systems.

Beachhead market

- Regulated enterprises already experimenting with AI agents
- Finance and procurement workflows where action risk is easy to quantify
- Organisations with formal risk, compliance and cybersecurity functions
- Departments willing to begin with a paid, contained pilot

Bottom-up opportunity model

Stage	Customer target	Average recurring value	Commercial objective
Year 1	6 customers	€50,000	Validate three paid pilots and secure initial contracts
Year 2	18 customers	€70,000	Expand across early regulated sectors
Year 3	45 customers	€90,000	Reach repeatable enterprise sales and break-even

Year 5	160 customers	€130,000	Scale across industries and larger deployments
--------	---------------	----------	--

These are internal planning assumptions, not externally validated market-size claims. The first commercial priority is to prove willingness to pay through paid pilots and measurable risk reduction.

17. Go-to-Market, Marketing and Main Partners

Growth plan

Stage 1 - Prototype: Demonstrate high-value payment, customer-data transfer and low-risk internal analysis scenarios.

Stage 2 - Paid pilots: Run controlled pilots in three regulated organisations and measure prevented risks, approval time and user trust.

Stage 3 - Sector packages: Develop dedicated banking, insurance, healthcare, government and technology policy templates.

Stage 4 - Partner-led scale: Expand through cybersecurity consultancies, systems integrators and enterprise AI platforms.

Marketing approach

- Direct enterprise sales to risk, compliance, security and AI-governance leaders
- Demonstrations at AI-governance, banking and cybersecurity events
- Compliance webinars and practical industry reports
- Pilot case studies showing measurable prevented risk and faster approval
- Strategic introductions through university, innovation and professional networks

Main partners

Partner type	Contribution
Pilot customers	Real workflows, policies, testing and commercial validation
AI platform providers	Agent integration and model connectivity
Cloud and identity providers	Secure deployment, access control and enterprise infrastructure
Cybersecurity consultancies	Customer implementation and channel sales
Systems integrators	Connections to ERP, CRM, payments and workflow systems
Legal, compliance and audit partners	Policy templates, governance advice and credibility
Universities and research partners	Responsible-AI evaluation, talent and innovation support

18. Technology Model

VetoPoint combines deterministic organisational rules with AI-supported contextual risk analysis. The prototype is a safe demonstration and will not execute real payments, send real messages, sign contracts or delete records.

1. AI agent proposes an action
2. VetoPoint verifies agent identity and authority
3. Company policies and hard rules are checked
4. AI analyses contextual risk and explains its reasoning
5. Decision engine returns Approve, Escalate or Veto
6. Human approval is requested when required
7. The action and final decision are recorded

Prototype technology stack

- Responsive web interface
- Secure backend functions
- Relational database
- AI model API with structured output
- Deterministic rules engine
- Role-based access control
- Tamper-evident audit logging
- Simulated enterprise integrations and fallback demo logic

Future integrations

- Enterprise AI-agent platforms
- Banking and payment systems
- CRM and ERP systems
- Human-resources platforms
- Email and communication tools
- Cloud infrastructure
- Identity and access-management systems

19. Responsible AI Framework

Data privacy

Process only information required to evaluate an action. Use encryption, data minimisation, role-based access, limited retention, private deployment options and secure audit records.

Bias and fairness

Review policies and recommendations to ensure that risk controls do not unfairly target individuals, departments, nationalities or protected groups. Test with diverse scenarios and monitor false positives.

Transparency

Show why an action was flagged, which rules applied, which risks were detected and whether each recommendation came from AI or a human.

Human oversight

Critical financial, legal, employment, healthcare, safety and privacy actions may not be automatically approved by AI.

Responsible use of information

Use information only for action-risk assessment, policy enforcement and authorised audit. Do not repurpose prompt or employee data for unrelated surveillance, advertising or model training without explicit permission.

Contestability

Allow users to challenge a recommendation, provide evidence and request another review.

Responsible monitoring

Evaluate proposed actions rather than scoring employee worth, loyalty or productivity.

Core principle

VetoPoint does not ask organisations to trust another AI blindly. It uses transparent rules, bounded AI analysis and human authority to decide when AI can be trusted to act.

20. Social, Ethical and Environmental Value

Value area	Contribution
Social value	Reduces the risk that automated actions harm customers, patients, employees or citizens and improves trust in responsible AI adoption.
Ethical value	Preserves human accountability, provides explanations, supports challenge and prevents critical decisions from being hidden inside an automated workflow.
Organisational value	Enables faster innovation without abandoning policy, compliance or oversight.
Environmental responsibility	Environmental benefit is not the primary claim. VetoPoint will minimise unnecessary processing through risk-based routing, data minimisation and appropriately sized models, and will report its infrastructure choices transparently.

21. Development Roadmap

Phase 1 - Competition prototype

- Landing page
- Decision Gate
- Three demonstration scenarios
- Risk-result screen
- Approval workflow
- Audit trail
- Executive dashboard
- Responsible-AI page

Phase 2 - Minimum viable product

- Authentication and role-based permissions
- Custom policies
- Secure database
- API integration
- Configurable approval workflows
- Multi-agent support
- Basic enterprise reporting

Phase 3 - Pilot-ready product

- Independent security testing
- Private deployment option
- Enterprise identity integration
- Customer-specific policies
- Monitoring and alerts
- Compliance reports
- Service-level support

Phase 4 - Scalable enterprise platform

- Multi-organisation architecture

- Advanced analytics
- Fraud-detection modules
- Agent identity verification
- Real-time integrations
- Industry-specific policy libraries
- Partner ecosystem

22. Illustrative Financial Projection

The following figures are internal planning assumptions, not guaranteed results.

Year	Customers	Estimated revenue	Estimated operating result
Year 1	6	€420,000	-€780,000
Year 2	18	€1.65 million	-€490,000
Year 3	45	€5.13 million	€910,000
Year 4	90	€12.02 million	€3.78 million
Year 5	160	€24.63 million	€9.69 million

Expected break-even

- Operational break-even during Year 3
- Approximately 35-40 customers
- €3.5-€4 million in recurring annual revenue
- Average annual customer value of approximately €90,000

Example medium-enterprise customer

- Annual subscription: €95,000
- Implementation: €35,000
- Premium module: €25,000
- First-year revenue: €155,000
- Following-year recurring revenue: €120,000

23. Main Costs

Product development

- Software engineering
- AI integration
- User-interface development
- Testing
- Cloud infrastructure

Security and compliance

- Penetration testing
- Legal review
- Data protection
- Security monitoring
- Insurance
- Compliance documentation

Sales and customer acquisition

- Enterprise sales team
- Demonstrations
- Industry events
- Pilot support
- Partnership development

Customer implementation

- System integration
- Policy configuration
- Staff training
- Technical support
- Customer success

24. Investment and Support Required

Initial investment sought
€750,000

Use of funds	Share
Product and AI development	40%
Cybersecurity and compliance	20%
Enterprise integrations	15%
Sales and pilot acquisition	15%
Legal and operations	10%

Non-financial support requested

- Access to three enterprise pilot customers
- Mentoring from cybersecurity, compliance and enterprise-sales experts
- Technical introductions to AI-agent and enterprise-platform providers
- Legal and responsible-AI guidance
- Visibility through investor, accelerator and industry networks

Twelve-month objectives

- Build a production-ready minimum viable product
- Complete independent security testing
- Launch three paid enterprise pilots
- Convert at least two pilots into annual contracts
- Build initial enterprise integrations
- Reach €300,000-€500,000 in contracted annual recurring revenue

25. Key Risks and Mitigation

Risk	Mitigation
Strong cybersecurity competitors	Focus on action-level governance, authorisation and audit rather than general prompt security.
Long enterprise sales cycles	Offer paid departmental pilots with clear success measures and contract conversion paths.

Integration complexity	Start with a small number of common agent workflows and a standard API.
Low trust in AI risk analysis	Use deterministic rules, transparent explanations and mandatory human approval for critical actions.
False warnings slow operations	Allow configurable thresholds, policy testing and continuous improvement.
Regulatory and legal uncertainty	Design for documentation, human authority, data minimisation and configurable regional policies.
Security of VetoPoint itself	Use secure development, independent testing, least-privilege access and private deployment options.

26. Success Indicators

- Number of AI actions reviewed
- Number and percentage of high-risk actions prevented
- Financial value of protected transactions
- Reduction in unauthorised actions
- Average approval time
- False-positive and false-negative rates
- Percentage of actions with complete audit evidence
- Pilot-to-contract conversion rate
- Annual recurring revenue
- Customer retention
- Number of protected AI agents
- Customer and approver satisfaction

27. Five-Minute Investor Pitch Structure

Keep each screen focused on one message. Aim for approximately 25-30 seconds per screen and reserve the final 20 seconds for the ask and closing line.

#	Pitch screen	Key message
1	Company name, logo and slogan	VetoPoint AI - Others secure the prompt. VetoPoint AI secures the action.
2	The problem	AI agents can transfer money, send data and change systems before organisations can intervene.
3	Target customer	Regulated enterprises, with banks and insurers as the first beachhead.
4	The AI solution	An independent control layer returning Approve, Escalate or Veto.
5	How it works	Agent identity -> policy rules -> AI risk analysis -> human approval -> audit record.
6	Business model	Annual subscriptions, paid pilots, implementation and premium modules.
7	Market opportunity	Start with regulated organisations deploying agents; expand through sector packages and partners.
8	Social, ethical or environmental value	Human authority, safer automation, transparency, privacy, fairness and efficient risk-based processing.
9	Team roles	Trinorte: Rubina Danilova - Marketing and Presentation Lead; Brais Muiño Otero - AI and Product Lead; Adrián Garrote Camoira - Team Leader and Business Strategist.
10	Investment or support required	€750,000, three paid pilot customers, security mentoring and integration partners.

Opening line

AI agents are beginning to act. Who checks their decisions before money, data or systems are affected?

Closing line

We are not asking businesses to trust another AI. We are giving them the infrastructure to decide when AI can be trusted to act.

28. AI Knowledge Quiz Readiness

The provided competition instructions mention an AI knowledge quiz but do not include the final list of topics. The team should confirm the official quiz scope with the organiser. Until then, prepare the following core areas:

- Difference between artificial intelligence, machine learning, generative AI and AI agents
- How prompts, models, tools and data sources work together
- Hallucination, uncertainty and the need for verification
- Bias and fairness in AI systems
- Privacy, security and responsible information use
- Transparency, explainability and human oversight
- Difference between deterministic rules and AI recommendations
- Benefits and limitations of AI in business
- Responsible AI principles used in VetoPoint
- The team's own product architecture, business model and ethical safeguards

Team preparation method

Each member should be able to explain VetoPoint in one minute, answer one technical question, one business question and one responsible-AI question, and distinguish clearly between what the prototype demonstrates and what a production system would require.

29. Competition Requirements Compliance Matrix

Requirement	How addressed	Location	Status
Competition theme	AI-enabled company solving a real business and societal problem	Sections 2, 5 and 6	Complete
Diverse team of three	Three named members, complementary roles and functional diversity statement	Section 4	Complete
Team name	Trinorte	Sections 1, 4 and 11	Complete
Company name and slogan	VetoPoint AI; Others secure the prompt...	Sections 1 and 11	Complete
Mission, vision, values, purpose	Dedicated identity section	Section 3	Complete
Problem and affected users	Problem, risks and stakeholders	Section 5	Complete
AI solution and how it works	Approve, Escalate, Veto workflow	Sections 6 and 18	Complete
Role of AI	Contextual analysis, risk scoring, explanation	Section 7	Complete
Improvement in speed/quality/cost/access/UX	Dedicated impact table	Section 7	Complete
Logo and brand colours	Logo, meaning and palette	Section 11	Complete

One-sentence description	Final wording included	Sections 1 and 11	Complete
Website content	All required sections, team details, contact email and CTA	Section 12	Complete; publish live URL
One-minute promotional video	Scene plan and support rationale	Section 13	Complete
Business model	Customers, income, costs, partners, marketing, growth	Sections 14-17 and 23	Complete
Five-minute pitch	Ten-screen investor structure	Section 27	Complete
Responsible AI	Privacy, fairness, transparency, oversight, responsible information use	Section 19	Complete
Social/ethical/environmental value	Dedicated value section	Section 20	Complete
Investment/support required	Funding and non-financial support	Section 24	Complete
AI knowledge quiz	Provisional preparation topics; official scope required	Section 28	Confirm with organiser

Final items to complete before submission

Publish the landing page and insert its live URL, and confirm the official AI Knowledge Quiz topics and competition date.

VetoPoint AI by Trinorte

Others secure the prompt. VetoPoint AI secures the action.