

VetoPoint AI

Competition Master Document

Team: **Trinorte** · Members: Brais, Adrián, Rina

Identity

Mission

To empower enterprises to adopt autonomous AI safely by providing an intelligent governance layer that prevents risky actions, ensures regulatory compliance, and keeps humans in control of critical decisions.

Vision

To become the global standard for AI action governance, enabling a future where businesses can fully trust autonomous AI to execute tasks without compromising security or ethics.

Main Purpose

To bridge the gap between AI innovation and enterprise security. VetoPoint exists to eliminate the fear of “rogue AI” or operational hallucinations, acting as an intelligent firewall so companies can automate tasks faster and safer.

Values

- **Security-First:** We prioritize enterprise data protection and risk mitigation above all else.
- **Transparency (Explainability):** We make AI decision-making visible, ensuring every AI action can be audited and understood.
- **Human-in-the-Loop:** We believe AI should augment human capabilities, not replace human judgment when the stakes are high.
- **Innovation Without Fear:** We build the guardrails so IT teams can confidently deploy new AI tools instead of blocking them.

Target Customer

- **Primary Buyers:** a cross-functional buying group — CISO, CRO (Chief Risk Officer), CCO, CIO and CFO — who jointly own financial control, cybersecurity, policy enforcement and accountable approval.
- **Company Profile:** regulated enterprises deploying autonomous AI agents, with an initial beachhead in banks and payment companies, insurance companies and government agencies, expanding into healthcare and other large enterprises.

01 · Problem — Why Now

AI adoption moved faster than governance. Today, 75% of knowledge workers already use AI at work, and 78% of AI users bring unapproved tools rather than company-approved ones (Microsoft 2024 Work Trend Index). But as agents move from *reading* to *writing*, the risk changes: a single hallucination can authorise a payout or leak data, leaving the company liable with no audit trail. Internal risk-scoring of unsupervised agent actions places this in the critical range (92/100).

IBM's 2025 Cost of a Data Breach Report adds that 63% of organisations have no AI-governance policy, 97% of AI-related breaches involved organisations lacking proper AI access controls, and heavy “shadow AI” raised average breach costs by approximately \$670,000. The EU AI Act turns this governance gap into a compliance deadline, not just a security concern.

Traditional cybersecurity stops at the prompt. **It doesn't stop the action.**

Anticipated Jury Question

Q: How is this different from a normal antivirus or DLP?

A: Traditional security guards the input — what employees type into AI. We govern the action coming out. Antivirus asks “is this file dangerous?”; we ask “should this AI do this right now?” and we decide before it executes.

02 · Solution — The Decision Gate

VetoPoint AI is the control layer between agents and systems. Every request goes through four steps:

Intercept → Recognise Intent → Score & Verify → Act & Log

Score & Verify combines AI risk-analysis with deterministic rules: deterministic rules are hard limits that always win; the AI adds context by reading intent, spotting unusual risk combinations, scoring the action and explaining why. The AI cannot override a hard rule or auto-approve a critical action — it advises, humans decide.

Three Outcomes

- **Approve:** low-risk actions clear instantly.
- **Veto:** actions that break policy are stopped before execution.
- **Escalate:** the action is held, a named human is alerted, and only upon approval does it resume. The agent waits — it doesn't crash. This is the state nobody else in the category ships.

Live demo example: an agent tries to wire €48,000 to a new supplier. The request is held; a named human is alerted; only on approval does the transfer proceed.

Anticipated Jury Question

Q: What exactly does the AI do vs. the rules? Isn't it just an LLM guessing?

A: Deterministic rules are hard limits — they always win. The AI adds context: reads intent, spots unusual risk combinations, scores it and explains why. The AI cannot override a hard rule or auto-approve a critical action. It advises; humans decide.

Improvement

- **Cost:** Prevents massive financial and reputational disasters caused by unsupervised agent actions and costly breaches.

- **Quality / Security:** Provides an immutable audit log of every decision — why an action was approved, escalated or vetoed.
- **Speed:** Enables safe automation to proceed without banning AI agents outright. Trust, transparency and human oversight, without breaking the agent's run.

03 · Market Opportunity — Beachhead

The timing is regulatory. The EU AI Act is phasing in now — transparency obligations from August 2026, high-risk system governance from December 2027 — and VetoPoint is built as the exact control layer this regulation demands. Our **Serviceable Obtainable Market is approximately €3.9M in ARR over a 3-year target**, focusing on EU enterprises — banks, insurers, government — that have just given an agent write-access and recognised the risk.

The same control layer answers different urgent risks across sectors. Each segment buys governance, but the trigger, buyer, measurable outcome and expansion path differ.

Segment 01 — Banking (Flagship Beachhead)

Banks and payment companies face immediate, quantifiable exposure when AI agents can initiate payments, alter records or approve transactions. The first beachhead is a controlled finance or procurement workflow.

Priority buyer: CRO · CISO · CCO · CFO — a cross-functional group owning financial control, cybersecurity, policy enforcement and accountable approval.

Land motion: €25K finance pilot — three months, one department, up to three AI agents and three approval workflows, with a measurable evaluation report.

Measurable value: reduce unmanaged high-risk AI actions before execution, without sending every low-risk action to manual review. Tracked via: high-risk actions prevented, financial value protected, average approval time, audit completeness.

Commercial package: €250K+ subscription, €75K–€150K implementation, plus optional transaction-control, fraud, identity and executive-analytics modules.

Expansion path: payments → procurement → operations — begin with one quantifiable workflow, add protected agents and modules, then extend the control layer across departments and systems.

Segment 02 — Insurance: claims, fraud signals and sensitive customer data.

Segment 03 — Government: citizen impact, procurement and public auditability.

Beyond the three beachhead segments, the addressable market extends to healthcare and other large enterprises deploying autonomous AI agents.

Land with a pilot. Scale through control. Renew as infrastructure. VetoPoint uses an Enterprise SaaS model, decoupled from compute costs. We land with a paid pilot, giving a department three months to prove value; proven pilots move to annual subscriptions priced on **governed write-paths, not per-agent** — so revenue scales with the client’s risk exposure while avoiding procurement friction when they consolidate agents. Commercial principle: services help the customer deploy; subscriptions and modules create the durable company.

Business Model Canvas

Foundation	Offer & Relationship	Market & Economics
Key Partners: pilot customers; AI platform providers; cloud/identity providers; cybersecurity consultants; systems integrators; legal/compliance/audit partners	Value Proposition: stop harmful actions before execution; approve low-risk actions automatically; escalate uncertainty to the right human; enforce policy consistently; create complete audit evidence	Customer Segments: regulated enterprises using AI agents — banks, insurers, government, healthcare and large enterprises. Buyers: CISO, CRO, CCO, CIO, CFO
Key Activities: build and secure the Decision Gate; develop policy and risk engines; integrate enterprise systems; deliver paid pilots; create sector policy packages	Customer Relationships: high-touch enterprise discovery; three-month paid pilot; policy co-design and onboarding; customer success and premium support; renewal and expansion reviews	Cost Structure: software and AI engineering; cybersecurity and compliance; enterprise integrations; sales and pilot acquisition; customer implementation and support
Key Resources: Decision Gate and policy IP; secure technical architecture; enterprise integration capability; pilot evidence and audit data; cross-functional founding team	Channels: direct enterprise sales; AI-governance and cybersecurity events; webinars and industry reports; pilot case studies; strategic introductions	Revenue Streams: annual enterprise subscriptions; three-month paid pilots; implementation and integration fees; usage charges for protected actions; premium compliance and identity modules

Revenue Engine

Paid Pilot — €25,000: three-month controlled deployment for one department (up to three AI agents, three approval workflows, risk analysis, policy enforcement, audit trail, dashboard, evaluation report). Creditable toward the first annual contract.

Plan	Annual Licence	Designed For
VetoPoint Core	€40,000	Small regulated companies or one department
VetoPoint Business	€95,000	Medium organisation protecting several AI agents
VetoPoint Enterprise	€250,000+	Banks, governments and large corporations

VetoPoint Business (€95,000/year) includes: multi-agent control; configurable policies and approvals; executive reporting; role-based governance; expansion-ready module architecture.

Implementation: None (pilot only) · Small €15,000 · Medium €35,000 · Enterprise €100,000 · Complex enterprise €150,000.

Customer value simulator — Business plan example: annual licence €95,000 + €35,000 implementation + one premium module €25,000.

- First-year contract value: **€155,000**
- Following-year recurring value: **€120,000**
- Three-year customer value: **€395,000**
- Recurring share, year one: **77%** — meets the $\geq 75\%$ mature revenue-quality objective

Premium modules (annual): financial transaction control €30,000 · privacy/data-governance €25,000 · AI Act compliance reporting €20,000 · advanced fraud detection €40,000 · agent identity verification €30,000 · executive risk analytics €25,000 · private deployment €50,000–€150,000.

Professional services: responsible-AI audit €15,000–€50,000 · custom policy design €10,000–€30,000 · training €5,000–€20,000 · security assessment €15,000–€40,000 · premium support 15–20% of annual licence.

The objective is for at least 75% of revenue to remain recurring, rather than relying too heavily on consulting.

Go-to-Market Motion

Prove risk reduction, package the learning, then scale through partners.

1. **Prototype:** demonstrate the control point — show VetoPoint can pause an AI-proposed action, apply policy, analyse risk, return Approve/Escalate/Veto, request human authority and record the decision, across three high-contrast scenarios (risky supplier payment, customer-data transfer, low-risk internal analysis).
2. **Paid Pilots:** convert category proof into commercial action via contained three-month deployments.
3. **Sector Packages:** turn repeatable policies into packaged offerings per vertical (banking, insurance, government).
4. **Partner Scale:** use systems integrators and consultancies to distribute and integrate at scale.

Primary channel at this stage: direct demos and competition, website, pitch and university/innovation networks.

“Others secure the prompt. VetoPoint secures the action.”

The category has consolidated quickly. But here is the gap: everyone can **“Allow”** or **“Block.”** Nobody ships the **“Review”** state on the write-path.

Substitutes such as AWS Cedar evaluate deterministic rules for a fraction of a cent. But Cedar is binary — it cannot pause an action and wait for human approval. VetoPoint costs €0.0003 per action, around 12× more, because it performs semantic evaluation and provides an approval interface that resumes the run once a human decides. We compete on preventing multi-million-dollar mistakes without breaking the workflow.

Anticipated Jury Question

Q: Why charge 12× more per action than a rules engine like AWS Cedar?

A: Cedar evaluates deterministic rules only — it can Allow or Block, nothing in between. VetoPoint adds semantic risk evaluation and the “Escalate” state: pausing an action for a named human and resuming it on approval. That extra layer is what prevents multi-million-dollar mistakes without breaking the agent’s workflow.

A bottom-up path to break-even in Year 3. These figures are internal planning assumptions — not externally validated market-size claims or guaranteed results.

Five-Year Projection

Year	Customers	Revenue	Operating Result	Commercial Objective
1	6	€420,000	–€780,000	Validate three paid pilots and secure initial contracts
2	18	€1.65 million	–€490,000	Expand across early regulated sectors
3	45	€5.13 million	€910,000	Reach repeatable enterprise sales and break-even
4	90	€12.02 million	€3.78 million	Scale deployments and partner channels
5	160	€24.63 million	€9.69 million	Expand across industries and larger deployments

The average contract value rises because the customer mix gradually shifts from pilots and departmental deployments toward enterprise-wide contracts.

How the projection is calculated — Year 1: 6 customers × €50,000 recurring revenue = €300,000; 6 implementations × €20,000 = €120,000. **Total: €420,000.** **Year 3:** 45 customers × €90,000 recurring revenue = €4.05 million; 27 new implementations × €25,000 = €675,000; modules and services = €405,000. **Total: €5.13 million.** **Year 5:** 160 customers × €130,000 recurring revenue = €20.8 million; 70 new implementations × €25,000 = €1.75 million; premium modules and services = €2.08 million. **Total: €24.63 million.**

Cost Structure

Direct operating costs are approximately 16–24% of revenue (cloud infrastructure, AI-model usage, security monitoring, support, storage, licences). Gross margin improves as the company scales:

Year	Est. Gross Margin
Year 1	76%
Year 2	79%
Year 3	82%
Year 4	83%
Year 5	84%

Main operating expenses: software engineering, cybersecurity/penetration testing, compliance and legal work, enterprise sales, customer success, product development, insurance, administration. The company is expected to operate at a loss initially given the development, certification and sales-cycle demands of enterprise software.

Break-Even Point

VetoPoint reaches operational break-even during **Year 3**. Simplified break-even target: approximately 35–40 customers, average contract value €85,000–€95,000, annual recurring revenue of roughly €3.5–€4 million.

Initial Investment Request

Investment sought: €750,000 — to turn the control concept into pilot-ready infrastructure.

Use of Funds	Share	Amount
Product and AI development	40%	€300,000
Cybersecurity and compliance	20%	€150,000
Enterprise integrations	15%	€112,500
Sales and pilot acquisition	15%	€112,500
Legal and operations	10%	€75,000

18-Month Proof Points

The business model succeeds only when pilots convert. Within 18 months, the investment should achieve:

1. Production-ready MVP — authentication, custom policies, secure database, approval workflows, multi-agent support, enterprise reporting.
2. Independent security testing — secure development, penetration testing, compliance documentation.
3. Three paid enterprise pilots — proving prevented risk, approval speed, audit evidence and customer trust.
4. Two pilot conversions — at least two paid pilots converted into annual subscriptions with a clear expansion pathway.
5. Initial integrations — with enterprise AI-agent platforms, identity systems, workflows and operational tools.
6. €300,000–€500,000 in contracted annual recurring revenue — establishing recurring commercial proof before broad partner-led scale.

Anticipated Jury Question

Q: Why raise €750,000, if Year-1 burn is €780,000?

A: The €750K funds the first 12–14 months alongside pilot revenue already being booked. Hitting that milestone de-risks a follow-on round ahead of break-even. The company raises lean by design.

Pitch-Ready Explanation

VetoPoint operates through annual enterprise subscriptions ranging from €40,000 to more than €250,000. Customers also pay for implementation, protected actions, private deployment and specialised compliance modules. Our target is to reach 45 institutional customers and more

than €5 million in annual revenue by Year 3, becoming profitable as recurring subscription revenue scales.

Strong Financial Statement

One protected action can prevent a loss worth millions. That allows VetoPoint to charge enterprise prices while still delivering a clear financial return to the customer.

Company and Brand Design

- **Company Name:** VetoPoint AI
- **Team Name:** Trinorte
- **Slogan:** *“Others secure the prompt. VetoPoint secures the action.”*
- **Logo & Brand Colours:** see brand assets submitted separately (logo.png, logo-static.svg)
- **Company Description:** VetoPoint AI is the enterprise control layer that evaluates, approves, escalates or vetoes consequential AI actions before execution — creating an independent decision point between an AI agent’s intention and real-world execution, so companies can automate faster without losing control.

Responsible AI

- **Data Privacy:** The product is itself a privacy control (strong GDPR angle). The EU AI Act, with broad application from 2 August 2026, pressures companies to document and govern AI use — VetoPoint is built for exactly that.
- **Bias / Fairness:** Risk-scoring models are trained and audited to avoid discriminating by department, seniority or user profile, and are reviewed periodically to detect drift or unintended bias in what gets flagged.
- **Transparency:** Every risk score and approve/escalate/veto decision is explainable and logged, so employees and auditors can see exactly why an action was flagged.
- **Human Oversight:** High-risk actions always require a human “Veto” decision; the system only automates what falls within pre-approved low-risk thresholds.
- **Responsible Use:** VetoPoint follows privacy-by-design principles itself, minimising data retention and processing only what is necessary to assess risk.

Suggested Team Roles

- **Adrián** — Team Leader and Business Strategist (drives strategy; presents Problem, Market Opportunity and Financials)
- **Brais** — AI and Product Lead (builds the product; presents Solution and Competition)
- **Rubina** — Marketing and Presentation Lead (leads market engagement; presents the Hook and Live Demo)

“We are Trinorte: Brais builds the product, Adrián drives the strategy, and Rubina leads market engagement. We are not asking businesses to trust another AI. We are giving them the infrastructure to decide when AI can be trusted to act.”

“Others secure the prompt. VetoPoint secures the action.”